

Class 12: Transcriptomics and the analysis of RNA-Seq data

(Suraj Sidhu: A18512793)

Table of contents

Background	1
Data Import	1
Toy differential gene expression	2
DESeq Analysis	8
Volcano Plot	10
Save our Results	11

Background

Today we will analyze some RNASeq data from Himes et al. on the effects of a common steroid (dexamethasone) on airway smooth muscle cells (ASM cells).

Our starting point is the “counts” dat and “metadata” that contain the count values for each gene in their different experiments.

Data Import

```
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")  
  
head(counts)
```

	SRR1039508	SRR1039509	SRR1039512	SRR1039513	SRR1039516
ENSG000000000003	723	486	904	445	1170
ENSG000000000005	0	0	0	0	0
ENSG000000000419	467	523	616	371	582
ENSG000000000457	347	258	364	237	318
ENSG000000000460	96	81	73	66	118
ENSG000000000938	0	0	1	0	2
	SRR1039517	SRR1039520	SRR1039521		
ENSG000000000003	1097	806	604		
ENSG000000000005	0	0	0		
ENSG000000000419	781	417	509		
ENSG000000000457	447	330	324		
ENSG000000000460	94	102	74		
ENSG000000000938	0	0	0		

Q1 How many genes are in this dataset ?

```
nrow(counts)
```

[1] 38694

Q2. How many ‘control’ cell lines do we have?

```
table(metadata$dex)
```

```
control treated
      4       4
```

Toy differential gene expression

To start our analysis let's calculate the mean counts for all genes in the control experiments.

1. Extract all “control” columns from the `counts` object.
2. Calculate the mean for all rows (i.e. genes) of these “control” columns 3-4. Do the same for “treated”
3. Compare these `control.mean` and `treated.mean` values.

Q3. How would you make the above code in either approach more robust? Is there a function that could help here?

```

control inds <- metadata$dex == "control"

control counts <- counts[ ,control inds]

control mean <- rowMeans( control counts )

head(control mean)

```

```

ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
    900.75          0.00        520.50        339.75        97.25
ENSG000000000938
    0.75

```

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called treated.mean)

```

treated inds <- metadata$dex == "treated"

treated counts <- counts[ ,treated inds]

treated mean <- rowMeans( treated counts )

head(treated mean)

```

```

ENSG000000000003 ENSG000000000005 ENSG000000000419 ENSG000000000457 ENSG000000000460
    658.00          0.00        546.00        316.50        78.75
ENSG000000000938
    0.00

```

Store these together for easy bookkeeping as meancounts

```


meancounts <- data.frame(control mean, treated mean)
head(meancounts)

```

	control mean	treated mean
ENSG000000000003	900.75	658.00
ENSG000000000005	0.00	0.00
ENSG000000000419	520.50	546.00
ENSG000000000457	339.75	316.50
ENSG000000000460	97.25	78.75
ENSG000000000938	0.75	0.00

Q5 (a). Create a scatter plot showing the mean of the treated samples against the mean of the control samples.


```
plot(meancounts)
```



Q5 (b). You could also use the ggplot2 package to make this figure producing the plot below. What geom_?() function would you use for this plot? We often talk to metrics like “log2 fold-change”

```
library(ggplot2)
meancounts <- data.frame(
  control = control.mean,
  treated = treated.mean
)

ggplot(meancounts, aes(x = control, y = treated)) +
  geom_point() +
  labs() +
  theme_minimal()
```



```
## treated/control
log2(10/20)
```


[1] -1

Q6. Try plotting both axes on a log scale. What is the argument to plot() that allows you to do this?

```
# log="xy" applies log scale to both x and y axes
plot(meancounts$control, meancounts$treated,
      xlab = "Control mean counts",
      ylab = "Treated mean counts",
      log = "xy")
```

Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted from logarithmic plot

Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted from logarithmic plot



Let's calculate the log2 fold change for our treated over our control mean counts.

```
meancounts$log2fc <- log2(meancounts[, "treated"] / meancounts[, "control"])
```

```
head(meancounts)
```

	control	treated	log2fc
ENSG000000000003	900.75	658.00	-0.45303916
ENSG000000000005	0.00	0.00	NaN
ENSG000000000419	520.50	546.00	0.06900279
ENSG000000000457	339.75	316.50	-0.10226805
ENSG000000000460	97.25	78.75	-0.30441833
ENSG000000000938	0.75	0.00	-Inf

A common “rule of thumb” is a log2 fold change cutoff of +2 and -2 to call genes “up regulated” or “down regulated”

Number of “up” genes

```
sum(meancounts$log2fc > +2, na.rm =T)
```

```
[1] 1846
```

Numbers of “down” genes

```
sum(meancounts$log2fc > -2, na.rm =T)  
  
[1] 22928  
  
zero.vals <- which(meancounts[,1:2]==0, arr.ind=TRUE)  
  
to.rm <- unique(zero.vals[,1])  
mycounts <- meancounts[-to.rm,]  
head(mycounts)
```

	control	treated	log2fc
ENSG00000000003	900.75	658.00	-0.45303916
ENSG00000000419	520.50	546.00	0.06900279
ENSG00000000457	339.75	316.50	-0.10226805
ENSG00000000460	97.25	78.75	-0.30441833
ENSG00000000971	5219.00	6687.50	0.35769358
ENSG00000001036	2327.00	1785.75	-0.38194109

Q7. What is the purpose of the arr.ind argument in the which() function call above? Why would we then take the first column of the output and need to call the unique() function?

arr.ind=TRUE makes which() return row and column indices of zeros. We take the first column to get the gene rows, and unique() ensures each gene is only counted once before removing them.

```
up.ind <- mycounts$log2fc > 2  
down.ind <- mycounts$log2fc < (-2)
```

Q8. Using the up.ind vector above can you determine how many up regulated genes we have at the greater than 2 fc level?

```
sum(up.ind, na.rm = TRUE)
```

[1] 250

Q9. Using the down.ind vector above can you determine how many down regulated genes we have at the greater than 2 fc level?

```
sum(down.ind, na.rm = TRUE)
```

```
[1] 367
```

Q10. Do you trust these results? Why or why not?

No not really because these numbers are just a rough exploratory estimate.

DESeq Analysis

Let's do this analysis properly and keep your inner stats nerd happy.

```
library(DESeq2)
```

```
Warning: package 'matrixStats' was built under R version 4.5.2
```

For DESeq analysis we need 3 things: 1. count values 2. metadata telling us about the columns in `countdata (colData)` 3. design of experiment

Our first function from DESeq2 will setup the input required for analysis by storing all these 3 things together.

```
dds <- dds <- DESeqDataSetFromMatrix(countData = counts,
                                         colData = metadata,
                                         design = ~dex)
```

```
converting counts to integer mode
```

```
Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in
design formula are characters, converting to factors
```

The main function of DESeq2 that runs the analysis is called `DESeq()`

```
dds <- DESeq(dds)
```

```
estimating size factors
```

```
estimating dispersions
```

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

```
results(dds)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 38694 rows and 6 columns

	baseMean	log2FoldChange	lfcSE	stat	pvalue
	<numeric>	<numeric>	<numeric>	<numeric>	<numeric>
ENSG00000000003	747.1942	-0.3507030	0.168246	-2.084470	0.0371175
ENSG00000000005	0.0000	NA	NA	NA	NA
ENSG000000000419	520.1342	0.2061078	0.101059	2.039475	0.0414026
ENSG000000000457	322.6648	0.0245269	0.145145	0.168982	0.8658106
ENSG000000000460	87.6826	-0.1471420	0.257007	-0.572521	0.5669691
...
ENSG00000283115	0.000000	NA	NA	NA	NA
ENSG00000283116	0.000000	NA	NA	NA	NA
ENSG00000283119	0.000000	NA	NA	NA	NA
ENSG00000283120	0.974916	-0.668258	1.69456	-0.394354	0.693319
ENSG00000283123	0.000000	NA	NA	NA	NA
	padj				
	<numeric>				
ENSG00000000003	0.163035				
ENSG00000000005	NA				
ENSG000000000419	0.176032				
ENSG000000000457	0.961694				
ENSG000000000460	0.815849				
...	...				
ENSG00000283115	NA				
ENSG00000283116	NA				
ENSG00000283119	NA				
ENSG00000283120	NA				
ENSG00000283123	NA				

```

res <- results(dds)
head(res)

log2 fold change (MLE): dex treated vs control
Wald test p-value: dex treated vs control
DataFrame with 6 rows and 6 columns
  baseMean log2FoldChange    lfcSE     stat   pvalue
  <numeric>      <numeric> <numeric> <numeric> <numeric>
ENSG000000000003 747.194195 -0.3507030  0.168246 -2.084470 0.0371175
ENSG000000000005  0.000000      NA       NA       NA       NA
ENSG000000000419 520.134160  0.2061078  0.101059  2.039475 0.0414026
ENSG000000000457 322.664844  0.0245269  0.145145  0.168982 0.8658106
ENSG000000000460  87.682625 -0.1471420  0.257007 -0.572521 0.5669691
ENSG000000000938  0.319167 -1.7322890  3.493601 -0.495846 0.6200029
  padj
  <numeric>
ENSG000000000003  0.163035
ENSG000000000005      NA
ENSG000000000419  0.176032
ENSG000000000457  0.961694
ENSG000000000460  0.815849
ENSG000000000938      NA

```


Volcano Plot

This is common summary result figure from these types of experiments and plot the log2 fold-change vs. the adjusted p-value.

```

plot(res$log2FoldChange, -log(res$padj))
abline(v=c(-2,2), col="red")

```



Save our Results

```
write.csv(res, file="my_results.csv")
```